

U-Net Lesion Segmentation

Useful Links

2D Unet

- <https://www.kaggle.com/keegil/keras-u-net-starter-lb-0-277>
(<https://www.kaggle.com/keegil/keras-u-net-starter-lb-0-277>)
- <https://github.com/jocicmarko/ultrasound-nerve-segmentation>
(<https://github.com/jocicmarko/ultrasound-nerve-segmentation>)
- <https://medium.com/@SeoJaeDuk/medical-image-segmentation-part-1-unet-convolutional-networks-with-interactive-code-d07231eb29bf> (<https://medium.com/@SeoJaeDuk/medical-image-segmentation-part-1-unet-convolutional-networks-with-interactive-code-d07231eb29bf>)
- <https://github.com/yihui-he/u-net> (<https://github.com/yihui-he/u-net>)
- <https://github.com/zhixuhao/unet> (<https://github.com/zhixuhao/unet>)

3D Unet

- Keras <https://github.com/ellisdg/3DUnetCNN> (<https://github.com/ellisdg/3DUnetCNN>)
- TF <https://github.com/zhengyang-wang/3D-Unet--Tensorflow> (<https://github.com/zhengyang-wang/3D-Unet--Tensorflow>)



```
In [2]: from keras.models import Model
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Input, Conv2D, Dropout, MaxPooling2D, Conv2DTranspose, Conv3D, MaxPooling3D, Conv3DTranspose, UpSampling3D
from keras.optimizers import Adam
from keras.callbacks import ModelCheckpoint
from keras.utils import to_categorical, Sequence
from keras import backend as K

from skimage.io import imread, imsave
from skimage.external.tiffio import TiffWriter

from scipy.ndimage import interpolation

from copy import deepcopy
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import glob
import time
```

C:\Users\Gerald\Anaconda3\envs\tf3\lib\site-packages\h5py__init__.py:34: FutureWarning: Conversion of the second argument of `issubdtype` from `float` to `np.float64` is deprecated. In future, it will be treated as `np.float64 == np.dtype(float).type`.

from ._conv import register_converters as _register_converters
Using TensorFlow backend.

1) 2-D U-Net with small images

I downsampled and cropped images to be 128 x 128.

Here I will start with 13 training images (only!) taken at evenly spaced intervals from the same volume (one brain). I will test on the rest of the slices of the volume

```
In [1]: x_train_path = 'data\\small2d\\train_img'
y_train_path = 'data\\small2d\\train_lesion'
x_valid_path = 'data\\small2d\\valid_img'
y_valid_path = 'data\\small2d\\valid_lesion'
x_test_path = 'data\\small2d\\test_img'
y_test_path = 'data\\small2d\\test_lesion'

img_size = (128, 64)
```

In [8]: *# make data generators*

```
# https://github.com/keras-team/keras/issues/3059#issuecomment-364787723
seed = 2018
data_gen_args = dict(
    rescale=1./255,
    rotation_range=10.,
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.1,
    shear_range=0.1,
    horizontal_flip=True,
    fill_mode='constant',
    cval=0)

# also consider deformation: https://gist.github.com/fmder/e28813c1e8721830ff9c
# https://www.kaggle.com/c/ultrasound-nerve-segmentation/discussion/22062

x_train_datagen = ImageDataGenerator(**data_gen_args)
y_train_datagen = ImageDataGenerator(**data_gen_args)
x_valid_datagen = ImageDataGenerator(rescale=1./255)
y_valid_datagen = ImageDataGenerator(rescale=1./255)
x_test_datagen = ImageDataGenerator(rescale=1./255)
y_test_datagen = ImageDataGenerator(rescale=1./255)

flow_args = dict(target_size=img_size,
                  color_mode='grayscale',
                  batch_size=32,
                  class_mode=None,
                  seed=seed)

x_train_generator = x_train_datagen.flow_from_directory(x_train_path, **flow_args)
y_train_generator = y_train_datagen.flow_from_directory(y_train_path, **flow_args)
train_generator = zip(x_train_generator, y_train_generator)

x_valid_generator = x_valid_datagen.flow_from_directory(x_valid_path, **flow_args)
y_valid_generator = y_valid_datagen.flow_from_directory(y_valid_path, **flow_args)
valid_generator = zip(x_valid_generator, y_valid_generator)

x_test_generator = x_test_datagen.flow_from_directory(x_test_path, **flow_args)
y_test_generator = y_test_datagen.flow_from_directory(y_test_path, **flow_args)
test_generator = zip(x_test_generator, y_test_generator)
```

Found 35 images belonging to 1 classes.

Found 35 images belonging to 1 classes.

Found 78 images belonging to 1 classes.

Found 78 images belonging to 1 classes.

Found 79 images belonging to 1 classes.

Found 79 images belonging to 1 classes.

```

In [5]: # from https://github.com/jocicmarko/ultrasound-nerve-segmentation/blob/master/tr
smooth = 1.

def dice(y_true, y_pred): # for numpy arrays
    y_true_f = y_true.reshape(-1)
    y_pred_f = y_pred.reshape(-1)
    intersection = np.sum(y_true_f*y_pred_f)
    if np.sum(y_true_f + y_pred_f)==0:
        return 1.0
    return (2. * intersection) / (np.sum(y_true_f) + np.sum(y_pred_f))

def dice_coef(y_true, y_pred):
    y_true_f = K.flatten(y_true)
    y_pred_f = K.flatten(y_pred)
    intersection = K.sum(y_true_f * y_pred_f)
    if K.sum(y_true_f + y_pred_f)==0:
        return 1.0
    return (2. * intersection + smooth) / (K.sum(y_true_f) + K.sum(y_pred_f) + sm

def dice_coef_loss(y_true, y_pred):
    return -dice_coef(y_true, y_pred)

```

```

In [6]: def plot_masks(data, labels, predictions=None):
plt.figure(figsize=(15, 8))
if predictions is None:
    predictions = np.zeros_like(labels)

for i, (img, mask, pred) in enumerate(zip(data, labels, predictions)):
    # image only
    plt.subplot(3,5,i+1)
    h = plt.imshow(np.squeeze(img), aspect='auto')
    h.set_cmap('gray')
    plt.axis('off')
    plt.title('%d%% of pixels' % (100*mask.mean()))

    # overlay mask + image
    plt.subplot(3,5,i+6)
    merged = 0.7*np.tile(img, (1,1,3))
    if pred.max()==0:
        merged[:, :, 1] += 0.7*np.squeeze(mask)
    else:
        merged[:, :, 0] += 0.7*np.squeeze(pred)
    merged[merged>1] = 1
    h = plt.imshow(merged, aspect='auto')
    plt.axis('off')
    if pred.max()==0:
        plt.title('True labels')
    else:
        plt.title('Predictions')

    # overlay mask + prediction
    plt.subplot(3,5,i+11)
    masks = np.concatenate([pred, np.zeros(mask.shape[:2]+(2,))], axis=2)
    masks[:, :, 1] += np.squeeze(mask)
    plt.imshow(masks, aspect='auto')
    plt.axis('off')
    if pred.max()>0:
        plt.title('Dice: %1.3f' % dice(mask, pred))


    if i==4:
        break

```

```
In [9]: # check dataset
for data_batch, labels_batch in train_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break

plot_masks(data_batch, labels_batch)
```

data batch shape: (32, 128, 64, 1)
labels batch shape: (32, 128, 64, 1)



```
In [10]: one_weight = 0.95 # class weights
zero_weight = 0.05

# consider pixel weights also:
# https://stackoverflow.com/questions/50255438/pixel-wise-loss-weight-for-image-seg

def weighted_binary_crossentropy(y_true, y_pred):

    # Original binary crossentropy (see losses.py):
    # K.mean(K.binary_crossentropy(y_true, y_pred), axis=-1)

    # Calculate the binary crossentropy
    b_ce = K.binary_crossentropy(y_true, y_pred)

    # Apply the weights
    weight_vector = y_true * one_weight + (1. - y_true) * zero_weight
    weighted_b_ce = weight_vector * b_ce

    # Return the mean error
    return K.mean(weighted_b_ce)
```

```

In [11]: # modified from https://www.kaggle.com/keegil/keras-u-net-starter-lb-0-277
def unet2d(loss_fn=dice_coef_loss, lr=1e-3, act='relu', init='he_normal', pad='same',
          inputs = Input(img_size + (1,)))

c1 = Conv2D(16, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c1 = Dropout(0.1) (c1) if dropout else c1
c1 = Conv2D(16, (3, 3), activation=act, kernel_initializer=init, padding=pad)
p1 = MaxPooling2D((2, 2)) (c1)

c2 = Conv2D(32, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c2 = Dropout(0.1) (c2) if dropout else c2
c2 = Conv2D(32, (3, 3), activation=act, kernel_initializer=init, padding=pad)
p2 = MaxPooling2D((2, 2)) (c2)

c3 = Conv2D(64, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c3 = Dropout(0.2) (c3) if dropout else c3
c3 = Conv2D(64, (3, 3), activation=act, kernel_initializer=init, padding=pad)
p3 = MaxPooling2D((2, 2)) (c3)

c4 = Conv2D(128, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c4 = Dropout(0.2) (c4) if dropout else c4
c4 = Conv2D(128, (3, 3), activation=act, kernel_initializer=init, padding=pad)
p4 = MaxPooling2D(pool_size=(2, 2)) (c4)

c5 = Conv2D(256, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c5 = Dropout(0.3) (c5) if dropout else c5
c5 = Conv2D(256, (3, 3), activation=act, kernel_initializer=init, padding=pad)

u6 = Conv2DTranspose(128, (2, 2), strides=(2, 2), padding=pad) (c5)
u6 = concatenate([u6, c4])
c6 = Conv2D(128, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c6 = Dropout(0.2) (c6) if dropout else c6
c6 = Conv2D(128, (3, 3), activation=act, kernel_initializer=init, padding=pad)

u7 = Conv2DTranspose(64, (2, 2), strides=(2, 2), padding=pad) (c6)
u7 = concatenate([u7, c3])
c7 = Conv2D(64, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c7 = Dropout(0.2) (c7) if dropout else c7
c7 = Conv2D(64, (3, 3), activation=act, kernel_initializer=init, padding=pad)

u8 = Conv2DTranspose(32, (2, 2), strides=(2, 2), padding=pad) (c7)
u8 = concatenate([u8, c2])
c8 = Conv2D(32, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c8 = Dropout(0.1) (c8) if dropout else c8
c8 = Conv2D(32, (3, 3), activation=act, kernel_initializer=init, padding=pad)

u9 = Conv2DTranspose(16, (2, 2), strides=(2, 2), padding=pad) (c8)
u9 = concatenate([u9, c1], axis=3)
c9 = Conv2D(16, (3, 3), activation=act, kernel_initializer=init, padding=pad)
c9 = Dropout(0.1) (c9) if dropout else c9
c9 = Conv2D(16, (3, 3), activation=act, kernel_initializer=init, padding=pad)

outputs = Conv2D(1, (1, 1), activation='sigmoid') (c9)

model = Model(inputs=[inputs], outputs=[outputs])
model.compile(optimizer=Adam(lr=lr), loss=loss_fn, metrics=[dice_coef])

```

```
model.summary()
```

```
return model
```



```
In [63]: K.clear_session()
model = unet2d(loss_fn=weighted_binary_crossentropy)
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	(None, 128, 64, 1)	0	
conv2d_1 (Conv2D)	(None, 128, 64, 16)	160	input_1[0][0]
conv2d_2 (Conv2D)	(None, 128, 64, 16)	2320	conv2d_1[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 64, 32, 16)	0	conv2d_2[0][0]
conv2d_3 (Conv2D) 1[0][0]	(None, 64, 32, 32)	4640	max_pooling2d_1[0][0]
conv2d_4 (Conv2D)	(None, 64, 32, 32)	9248	conv2d_3[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 32, 16, 32)	0	conv2d_4[0][0]
conv2d_5 (Conv2D) 2[0][0]	(None, 32, 16, 64)	18496	max_pooling2d_2[0][0]
conv2d_6 (Conv2D)	(None, 32, 16, 64)	36928	conv2d_5[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 16, 8, 64)	0	conv2d_6[0][0]
conv2d_7 (Conv2D) 3[0][0]	(None, 16, 8, 128)	73856	max_pooling2d_3[0][0]
conv2d_8 (Conv2D)	(None, 16, 8, 128)	147584	conv2d_7[0][0]

max_pooling2d_4 (MaxPooling2D)	(None, 8, 4, 128)	0	conv2d_8[0][0]
conv2d_9 (Conv2D)	(None, 8, 4, 256)	295168	max_pooling2d_4[0][0]
conv2d_10 (Conv2D)	(None, 8, 4, 256)	590080	conv2d_9[0][0]
conv2d_transpose_1 (Conv2DTrans	(None, 16, 8, 128)	131200	conv2d_10[0][0]
concatenate_1 (Concatenate)	(None, 16, 8, 256)	0	conv2d_transpo se_1[0][0]
conv2d_11 (Conv2D)	(None, 16, 8, 128)	295040	concatenate_1 [0][0]
conv2d_12 (Conv2D)	(None, 16, 8, 128)	147584	conv2d_11[0][0]
conv2d_transpose_2 (Conv2DTrans	(None, 32, 16, 64)	32832	conv2d_12[0][0]
concatenate_2 (Concatenate)	(None, 32, 16, 128)	0	conv2d_transpo se_2[0][0]
conv2d_13 (Conv2D)	(None, 32, 16, 64)	73792	concatenate_2 [0][0]
conv2d_14 (Conv2D)	(None, 32, 16, 64)	36928	conv2d_13[0][0]
conv2d_transpose_3 (Conv2DTrans	(None, 64, 32, 32)	8224	conv2d_14[0][0]
concatenate_3 (Concatenate)	(None, 64, 32, 64)	0	conv2d_transpo se_3[0][0]
			conv2d_4[0][0]

conv2d_15 (Conv2D) [0][0]	(None, 64, 32, 32)	18464	concatenate_3
conv2d_16 (Conv2D) [0]	(None, 64, 32, 32)	9248	conv2d_15[0]
conv2d_transpose_4 (Conv2DTrans [0]	(None, 128, 64, 16)	2064	conv2d_16[0]
concatenate_4 (Concatenate) se_4[0][0]	(None, 128, 64, 32)	0	conv2d_transpo conv2d_2[0][0]
conv2d_17 (Conv2D) [0][0]	(None, 128, 64, 16)	4624	concatenate_4
conv2d_18 (Conv2D) [0]	(None, 128, 64, 16)	2320	conv2d_17[0]
conv2d_19 (Conv2D) [0]	(None, 128, 64, 1)	17	conv2d_18[0]
=====			
=====			
Total params: 1,940,817			
Trainable params: 1,940,817			
Non-trainable params: 0			

```
In [64]: # checkpoint = ModelCheckpoint(filepath='weights.h5')
checkpoint = ModelCheckpoint(filepath='weights2.h5', monitor='val_loss', save_best_only=True)

history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=20,
    validation_data=valid_generator,
    validation_steps=100,
    verbose=1,
    callbacks=[checkpoint])
```

```
Epoch 1/20
100/100 [=====] - 52s 523ms/step - loss: 0.0475 - dice_coef: 0.1023 - val_loss: 0.0289 - val_dice_coef: 0.1578
Epoch 2/20
100/100 [=====] - 49s 495ms/step - loss: 0.0333 - dice_coef: 0.2439 - val_loss: 0.0136 - val_dice_coef: 0.4260
Epoch 3/20
100/100 [=====] - 36s 364ms/step - loss: 0.0212 - dice_coef: 0.3931 - val_loss: 0.0116 - val_dice_coef: 0.5929
Epoch 4/20
100/100 [=====] - 51s 507ms/step - loss: 0.0206 - dice_coef: 0.4165 - val_loss: 0.0101 - val_dice_coef: 0.5812
Epoch 5/20
100/100 [=====] - 64s 637ms/step - loss: 0.0197 - dice_coef: 0.4530 - val_loss: 0.0101 - val_dice_coef: 0.5694
Epoch 6/20
100/100 [=====] - 40s 398ms/step - loss: 0.0163 - dice_coef: 0.5278 - val_loss: 0.0102 - val_dice_coef: 0.6399
Epoch 7/20
100/100 [=====] - 49s 492ms/step - loss: 0.0184 - dice_coef: 0.4866 - val_loss: 0.0097 - val_dice_coef: 0.5990
Epoch 8/20
100/100 [=====] - 49s 491ms/step - loss: 0.0188 - dice_coef: 0.4835 - val_loss: 0.0097 - val_dice_coef: 0.6427
Epoch 9/20
100/100 [=====] - 40s 403ms/step - loss: 0.0189 - dice_coef: 0.4823 - val_loss: 0.0098 - val_dice_coef: 0.5993
Epoch 10/20
100/100 [=====] - 51s 505ms/step - loss: 0.0197 - dice_coef: 0.4618 - val_loss: 0.0093 - val_dice_coef: 0.6347
Epoch 11/20
100/100 [=====] - 40s 402ms/step - loss: 0.0176 - dice_coef: 0.4763 - val_loss: 0.0122 - val_dice_coef: 0.6471
Epoch 12/20
100/100 [=====] - 34s 339ms/step - loss: 0.0197 - dice_coef: 0.4674 - val_loss: 0.0098 - val_dice_coef: 0.6827
Epoch 13/20
100/100 [=====] - 44s 436ms/step - loss: 0.0178 - dice_coef: 0.5026 - val_loss: 0.0090 - val_dice_coef: 0.7077
Epoch 14/20
100/100 [=====] - 146s 1s/step - loss: 0.0181 - dice_coef: 0.4888 - val_loss: 0.0118 - val_dice_coef: 0.6341
Epoch 15/20
100/100 [=====] - 113s 1s/step - loss: 0.0173 - dice_coef: 0.4888 - val_loss: 0.0118 - val_dice_coef: 0.6341
```

```

oef: 0.4773 - val_loss: 0.0084 - val_dice_coef: 0.6931
Epoch 16/20
100/100 [=====] - 106s 1s/step - loss: 0.0196 - dice_c
oef: 0.4671 - val_loss: 0.0096 - val_dice_coef: 0.6664
Epoch 17/20
100/100 [=====] - 96s 964ms/step - loss: 0.0179 - dice
_coef: 0.5012 - val_loss: 0.0091 - val_dice_coef: 0.6638
Epoch 18/20
100/100 [=====] - 92s 916ms/step - loss: 0.0173 - dice
_coef: 0.5194 - val_loss: 0.0087 - val_dice_coef: 0.7010
Epoch 19/20
100/100 [=====] - 42s 416ms/step - loss: 0.0169 - dice
_coef: 0.4929 - val_loss: 0.0089 - val_dice_coef: 0.6736
Epoch 20/20
100/100 [=====] - 42s 420ms/step - loss: 0.0214 - dice
_coef: 0.4560 - val_loss: 0.0091 - val_dice_coef: 0.5957

```

```

In [43]: def plot_performance(history):
    dicec = history.history['dice_coef']
    val_dice = history.history['val_dice_coef']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(dicec))
    plt.figure(figsize=(10, 4))
    plt.subplot(1,2,1)
    plt.plot(epochs, dicec, label='Training dice')
    plt.plot(epochs, val_dice, label='Validation dice')
    plt.xlabel("Num of Epochs")
    plt.ylabel("dice")
    plt.title('Training and validation dice')
    plt.legend()

    plt.subplot(1,2,2)

    plt.plot(epochs, loss, label='Training loss')
    plt.plot(epochs, val_loss, label='Validation loss')
    plt.xlabel("Num of Epochs")
    plt.ylabel("Loss")
    plt.title('Training and validation loss')
    plt.legend()

    plt.show()

```

```
In [65]: plot_performance(history)
```



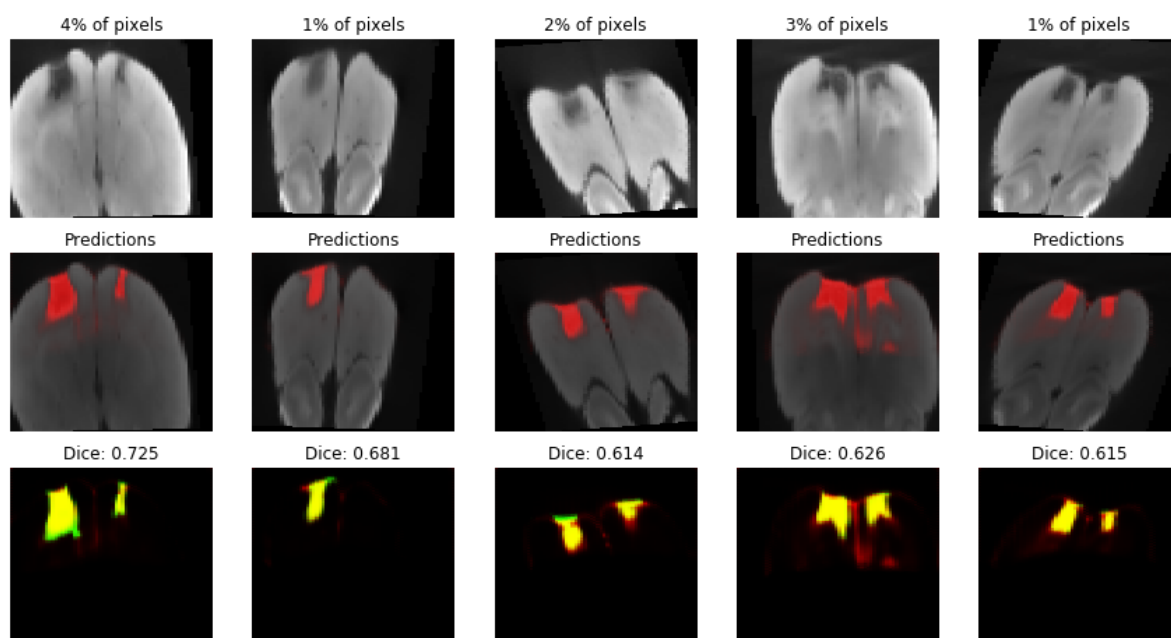
```
In [ ]: # Load best model
model.load_weights('weights2.h5')
```

```
In [25]: # training set performance
for data_batch, labels_batch in train_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break

predict_batch = model.predict_on_batch(data_batch)

plot_masks(data_batch, labels_batch, predict_batch)
```

data batch shape: (32, 128, 64, 1)
labels batch shape: (32, 128, 64, 1)

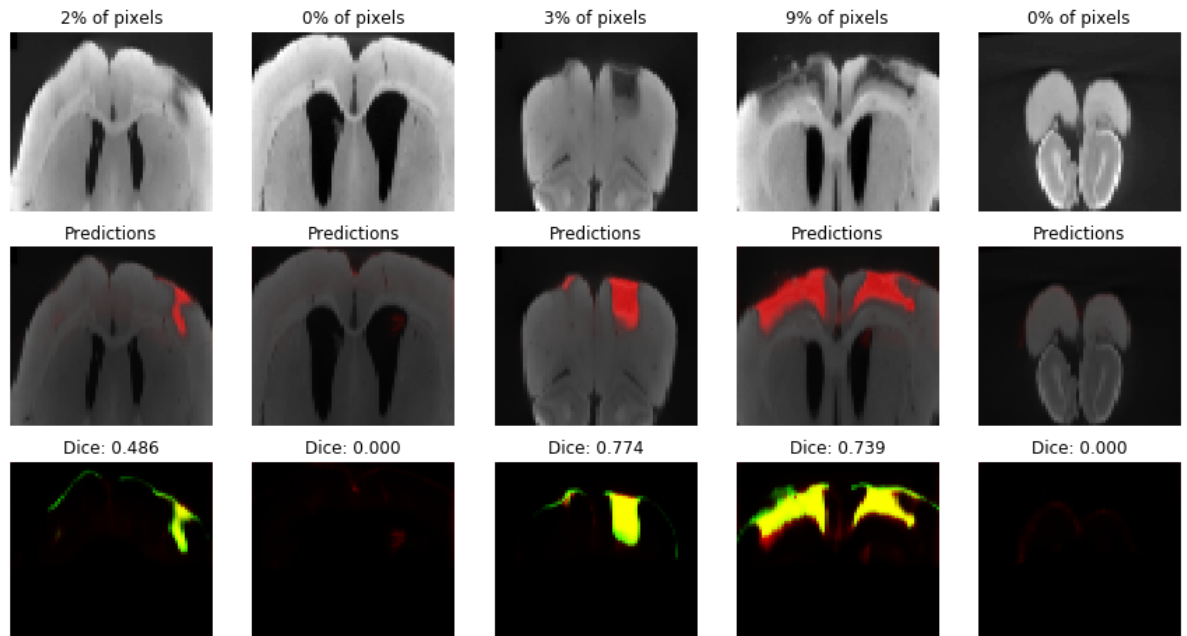


```
In [18]: # validation set performance
for data_batch, labels_batch in valid_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break

predict_batch = model.predict_on_batch(data_batch)

plot_masks(data_batch, labels_batch, predict_batch)
```

data batch shape: (32, 128, 64, 1)
labels batch shape: (32, 128, 64, 1)

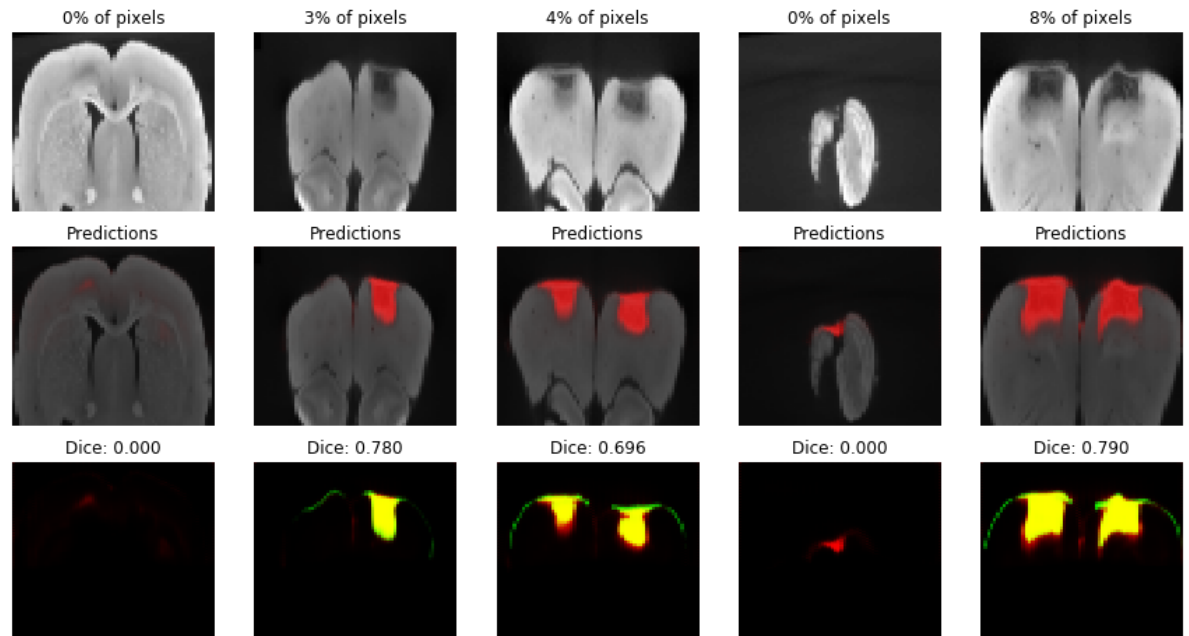


```
In [31]: # test set performance
for data_batch, labels_batch in test_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break

predict_batch = model.predict_on_batch(data_batch)

plot_masks(data_batch, labels_batch, predict_batch)
```

data batch shape: (32, 128, 64, 1)
labels batch shape: (32, 128, 64, 1)



3D U-Net


```

In [32]: vol_folder = 'data\\small3d\\'

vols = ['GP1825', 'GP1826', 'GP1827']

for (i, vol) in enumerate(vols):
    train_vol = imread(os.path.join(vol_folder, vol) + '.tif')
    train_vol = train_vol.astype('float') / train_vol.max()
    train_labels = imread(os.path.join(vol_folder, vol) + '_lesion.tif')

    factor = train_labels[train_labels > 0].min()
    train_labels = np.round(train_labels / factor).astype('uint8')

    if i == 0:
        X_train = np.moveaxis(np.expand_dims(np.expand_dims(train_vol, axis=3), axis=0), axis=-1, axis=1)
        y_train = np.moveaxis(np.expand_dims(np.expand_dims(train_labels, axis=3), axis=0), axis=-1, axis=1)
    else:
        X_train = np.append(X_train, np.moveaxis(np.expand_dims(np.expand_dims(train_vol, axis=3), axis=0), axis=-1, axis=1), axis=0)
        y_train = np.append(y_train, np.moveaxis(np.expand_dims(np.expand_dims(train_labels, axis=3), axis=0), axis=-1, axis=1), axis=0)

```

```

In [33]: # since Keras' ImageDataGenerator doesn't support 3D data, we make our own genera
# https://stanford.edu/~shervine/blog/keras-how-to-generate-data-on-the-fly

class VolumeDataGenerator(Sequence):
    def __init__(self, filenames, root, label_suffix='lesion', batch_size=1, dim=
        rotation_range=10, zoom_range=[0.9, 1.1], brightness_range=[0.8,
self.dim = dim
self.batch_size = batch_size
self.n_channels = n_channels
self.filenames = filenames
self.root = root
if self.batch_size > len(self.filenames):
    self.batch_size = len(self.filenames)
self.label = label_suffix
self.shuffle = shuffle
self.reset_index()

    # augmentation parameters
self.rotation_range = rotation_range
self.zoom_range = zoom_range
self.brightness_range = brightness_range
self.flip_axes = np.roll(flip_axes, -1) # y, z, x --> x, y, z
self.shift_range = shift_range

    def __len__(self):
        # number of batches per epoch
        return int(np.floor(len(self.filenames) / self.batch_size))

    def __getitem__(self, index):
        # Generate one batch of data
        indexes = self.indexes[index*self.batch_size:(index+1)*self.batch_size]

        files = [self.filenames[k] for k in indexes]

        X, y = self.load_files(files)

        for i, (Xi, yi) in enumerate(zip(X, y)):
            Xi, yi = self.augmentit(np.squeeze(Xi), np.squeeze(yi))
            X[i] = np.expand_dims(Xi, axis=3)
            y[i] = np.expand_dims(yi, axis=3)

        return X, y

    def reset_index(self):
        # Updates indexes after each epoch
        self.indexes = np.arange(len(self.filenames))
        if self.shuffle == True:
            np.random.shuffle(self.indexes)

    def load_files(self, files):

        X = np.empty((self.batch_size, *self.dim, self.n_channels))
        y = np.empty((self.batch_size, *self.dim, self.n_channels), dtype=int)

        for i, file in enumerate(files):

```

```

img = imread(os.path.join(self.root, file)+'.tif')
img = img.astype('float') / img.max()
X[i] = np.expand_dims(img, axis=3)

label = imread(os.path.join(self.root, file)+'_'+self.label+'.tif')
factor = label[label>0].min()
label = np.round(label/factor).astype('uint8')
y[i] = np.expand_dims(label, axis=3)

return X, y

## the following augmentation functions are copied from
# https://mlnotebook.github.io/post/dataaug/

def augmentit(self, thisim, thisseg):
    dims = self.dim
    aug = False

    if self.rotation_range>0:
        theta = np.random.uniform(-self.rotation_range,self.rotation_range)
        thisim = self.rotateit(thisim, theta)
        thisseg = self.rotateit(thisseg, theta, isseg=True)
        aug = True

    if not np.equal(self.zoom_range, 1).all():
        scalefactor = float(np.around(np.random.uniform(self.zoom_range[0],
        thisim = self.scaleit(thisim, scalefactor)
        thisseg = self.scaleit(thisseg, scalefactor, isseg=True)
        aug = True

    if not np.equal(self.brightness_range, 1).all():
        factor = float(np.around(np.random.uniform(self.brightness_range[0],
        thisim = self.intensifyit(thisim, factor)
        aug = True
        #no intensity change on segmentation

    if not np.equal(self.flip_axes, 0).all():
        thisim = self.flipit(thisim, self.flip_axes)
        thisseg = self.flipit(thisseg, self.flip_axes)
        aug = True

    if self.shift_range>0:
        offset = list(np.random.randint(-self.shift_range,self.shift_range,
        thisim = self.translateit(thisim, offset)
        thisseg = self.translateit(thisseg, offset, isseg=True)
        aug = True

    if aug:
        thisim, thisseg = self.cropit(thisim, thisseg)
        thisseg = self.resampleit(thisseg, dims, isseg=True)
        thisim = self.resampleit(thisim, dims)

    return thisim, thisseg

def scaleit(self, image, factor, isseg=False):
    order = 0 if isseg == True else 3

```

```

height, width, depth = image.shape
zheight              = int(np.round(factor * height))
zwidth               = int(np.round(factor * width))

if factor < 1.0:
    newimg = np.zeros_like(image)
    row    = int((height - zheight) / 2)
    col    = int((width - zwidth) / 2)
    newimg[row:row+zheight, col:col+zwidth, :] = interpolation.zoom(image

    return newimg

elif factor > 1.0:
    row    = (zheight - height) // 2
    col    = (zwidth - width) // 2

    newimg = interpolation.zoom(image[row:row+zheight, col:col+zwidth], (

    extrah = (newimg.shape[0] - height) // 2
    extraw = (newimg.shape[1] - width) // 2
    newimg = newimg[extrah:extrah+height, extraw:extraw+width, :]

    return newimg

else:
    return image

def resampleit(self, image, dims, isseg=False):
    order = 0 if isseg == True else 5

    image = interpolation.zoom(image, np.array(dims)/np.array(image.shape, dt

    if isseg:
        image[np.where(image==4)]=3

    return image if isseg else (image-image.min())/(image.max()-image.min())

def translateit(self, image, offset, isseg=False, mode='nearest'):
    order = 0 if isseg else 5
    mode = 'nearest'
    offset = (int(offset[0]), int(offset[1]), int(offset[2]))

    return interpolation.shift(image, offset , order=order, mode=mode)

def rotateit(self, image, theta, isseg=False):
    # modified using https://stackoverflow.com/questions/43922198/how-to-rotate-image
    order = 0 if isseg == True else 5
    for i, angle in enumerate(theta):
        image = interpolation.rotate(image, angle, mode='nearest', axes=(int(

    return image

def flipit(self, image, axes):

```

```

    if axes[0]:
        image = np.fliplr(image)
    if axes[1]:
        image = np.flipud(image)
    if axes[2]:
        image = np.flip(image, axis=2)

    return image

def intensifyit(self, image, factor):

    return image*float(factor)

def cropit(self, image, seg=None, margin=5):

    shortaxis = np.argmin(image.shape[:2])
    trimaxes = 0 if shortaxis == 1 else 1

    trim = image.shape[shortaxis]
    center = image.shape[trimaxes] // 2
    lrcenter = image.shape[shortaxis] // 2

    if seg is not None:

        hits = np.where(seg!=0)
        mins = np.amin(hits, axis=1)
        maxs = np.amax(hits, axis=1)
        segtrim = max(maxs-mins) + margin

        trim = segtrim
        center = np.mean(hits, 1, dtype=int)[0]
        lrcenter = np.mean(hits, 1, dtype=int)[1]

        if center - (trim // 2) > mins[0]:
            while center - (trim // 2) > mins[0]:
                center = center - 1
            center = center

        if center + (trim // 2) < maxs[0]:
            while center + (trim // 2) < maxs[0]:
                center = center + 1
            center = center

        if lrcenter - (trim // 2) > mins[1]:
            while lrcenter - (trim // 2) > mins[1]:
                lrcenter = lrcenter - 1
            lrcenter = lrcenter

        if lrcenter + (trim // 2) < maxs[1]:
            while lrcenter + (trim // 2) < maxs[1]:
                lrcenter = lrcenter + 1
            lrcenter = lrcenter

    top = max(0, center - (trim // 2) - margin//2)
    bottom = trim + margin if top == 0 else top + trim + (margin//2)

```

```
left = max(0, lrcenter - (trim//2) - margin//2)
right = trim + margin if left == 0 else left + trim + (margin//2)

# image[center-5:center+5, lrcenter-5:lrcenter+5, :] = 255
# image[top:bottom, left-2:left+2, :] = 255
# image[top:bottom, right-2:right+2, :] = 255
# image[top-2:top+2, left:right, :] = 255
# image[bottom-2:bottom+2, left:right, :] = 255

if bottom > image.shape[trimaxes]:
    bottom = image.shape[trimaxes]
    top = bottom - trim

if right > image.shape[shortaxis]:
    right = image.shape[shortaxis]
    left = right - trim

image = image[top: bottom, left:right]

if seg is not None:
    seg = seg[top: bottom, left:right]

    return image, seg
else:
    return image
```

```
In [34]: filenames = ['GP1825', 'GP1826', 'GP1827']
root = 'data\\small3d\\'

vol_size = (64, 64, 128)
train_gen = VolumeDataGenerator(filenames, root, rotation_range=10, zoom_range=[0,
                                     flip_axes=[1,0,1], shift_range=5)

noaug_gen = VolumeDataGenerator(filenames, root, rotation_range=0, zoom_range=[1,
                                     flip_axes=[0,0,0], shift_range=0, shuffle=False,

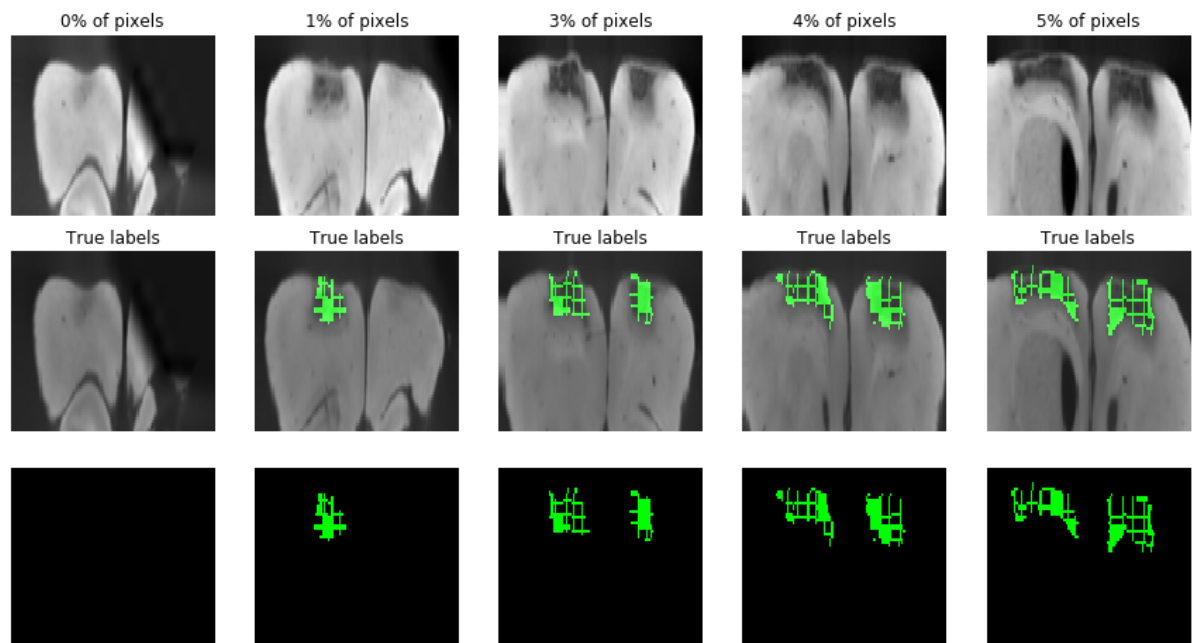
start = time.time()
X, y = train_gen.__getitem__(0)
print(time.time()-start)

# preview the augmentation
plot_masks(X[0][4:-1:8], y[0][4:-1:8]>1)
```

C:\Users\Gerald\Anaconda3\envs\tf3\lib\site-packages\scipy\ndimage\interpolation.py:600: UserWarning: From scipy 0.13.0, the output shape of zoom() is calculated with round() instead of int() - for these inputs the size of the returned array has changed.

"the returned array has changed.", UserWarning)

1.9856905937194824



The masks look like checkboards because I only sparsely annotated the volumes in each orthogonal dimension -- they are not perfect checkboards because of the random rotation.

```
In [36]: n_classes = y.max()+1
class_weights = np.array([0, 0.05, 0.95])

def weighted_categorical_crossentropy(y_true, y_pred):

    # categorical cross entropy
    y_true_fixed = K.relu(y_true-1) # shift category labels
    c_ce = K.sparse_categorical_crossentropy(y_true_fixed, y_pred)

    # Apply class weights
    weight_vector = K.squeeze(K.sum(K.one_hot(K.cast(y_true, 'int32'), n_classes)
    #   weight_vector = y_true * one_weight + (1. - y_true) * zero_weight
    weighted_c_ce = weight_vector * c_ce
    #   weighted_c_ce = c_ce
    #   weighted_c_ce = c_ce*class_weights

    # Return the mean error
    return K.mean(weighted_c_ce)
```



```

In [37]: # modified from https://www.kaggle.com/keegil/keras-u-net-starter-lb-0-277
def unet3d(loss_fn=weighted_categorical_crossentropy, lr=1e-3, act='relu', init='
    inputs = Input(vol_size + (1,))

    c1 = Conv3D(16, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c1 = Dropout(0.1) (c1) if dropout else c1
    c1 = Conv3D(16, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    p1 = MaxPooling3D((2, 2, 2)) (c1)

    c2 = Conv3D(32, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c2 = Dropout(0.1) (c2) if dropout else c2
    c2 = Conv3D(32, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    p2 = MaxPooling3D((2, 2, 2)) (c2)

    c3 = Conv3D(64, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c3 = Dropout(0.2) (c3) if dropout else c3
    c3 = Conv3D(64, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    p3 = MaxPooling3D((2, 2, 2)) (c3)

    c4 = Conv3D(128, (3, 3, 3), activation=act, kernel_initializer=init, padding=
    c4 = Dropout(0.2) (c4) if dropout else c4
    c4 = Conv3D(128, (3, 3, 3), activation=act, kernel_initializer=init, padding=

    # remove one layer of depth to save on model size / memory
    p4 = MaxPooling3D(pool_size=(2, 2, 2)) (c4)

    c5 = Conv3D(256, (3, 3, 3), activation=act, kernel_initializer=init, padding=
    c5 = Dropout(0.3) (c5) if dropout else c5
    c5 = Conv3D(256, (3, 3, 3), activation=act, kernel_initializer=init, padding=

    u6 = Conv3DTranspose(128, (2, 2, 2), strides=(2, 2, 2), padding=pad) (c5)
    u6 = concatenate([u6, c4])
    c6 = Conv3D(128, (3, 3, 3), activation=act, kernel_initializer=init, padding=
    c6 = Dropout(0.2) (c6) if dropout else c6
    c6 = Conv3D(128, (3, 3, 3), activation=act, kernel_initializer=init, padding=
    # end remove

    u7 = Conv3DTranspose(64, (2, 2, 2), strides=(2, 2, 2), padding=pad) (c6)
    u7 = concatenate([u7, c3])
    c7 = Conv3D(64, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c7 = Dropout(0.2) (c7) if dropout else c7
    c7 = Conv3D(64, (3, 3, 3), activation=act, kernel_initializer=init, padding=p

    u8 = Conv3DTranspose(32, (2, 2, 2), strides=(2, 2, 2), padding=pad) (c7)
    u8 = concatenate([u8, c2])
    c8 = Conv3D(32, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c8 = Dropout(0.1) (c8) if dropout else c8
    c8 = Conv3D(32, (3, 3, 3), activation=act, kernel_initializer=init, padding=p

    u9 = Conv3DTranspose(16, (2, 2, 2), strides=(2, 2, 2), padding=pad) (c8)
    u9 = concatenate([u9, c1])
    c9 = Conv3D(16, (3, 3, 3), activation=act, kernel_initializer=init, padding=p
    c9 = Dropout(0.1) (c9) if dropout else c9
    c9 = Conv3D(16, (3, 3, 3), activation=act, kernel_initializer=init, padding=p

    outputs = Conv3D(n_classes-1, (1, 1, 1), activation='softmax') (c9)

```

```
model = Model(inputs=[inputs], outputs=[outputs])  
model.compile(optimizer=Adam(lr=lr), loss=loss_fn, metrics=[loss_fn])  
model.summary()  
  
return model
```

```
In [38]: K.clear_session()
model = unet3d(loss_fn=weighted_categorical_crossentropy, dropout=False)
```

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	(None, 64, 64, 128, 0		
conv3d_1 (Conv3D)	(None, 64, 64, 128, 448		input_1[0][0]
conv3d_2 (Conv3D)	(None, 64, 64, 128, 6928		conv3d_1[0][0]
max_pooling3d_1 (MaxPooling3D)	(None, 32, 32, 64, 1 0		conv3d_2[0][0]
conv3d_3 (Conv3D)	(None, 32, 32, 64, 3 13856		max_pooling3d_1[0][0]
conv3d_4 (Conv3D)	(None, 32, 32, 64, 3 27680		conv3d_3[0][0]
max_pooling3d_2 (MaxPooling3D)	(None, 16, 16, 32, 3 0		conv3d_4[0][0]
conv3d_5 (Conv3D)	(None, 16, 16, 32, 6 55360		max_pooling3d_2[0][0]
conv3d_6 (Conv3D)	(None, 16, 16, 32, 6 110656		conv3d_5[0][0]
max_pooling3d_3 (MaxPooling3D)	(None, 8, 8, 16, 64) 0		conv3d_6[0][0]
conv3d_7 (Conv3D)	(None, 8, 8, 16, 128 221312		max_pooling3d_3[0][0]
conv3d_8 (Conv3D)	(None, 8, 8, 16, 128 442496		conv3d_7[0][0]

max_pooling3d_4 (MaxPooling3D)	(None, 4, 4, 8, 128) 0	conv3d_8[0][0]
conv3d_9 (Conv3D)	(None, 4, 4, 8, 256) 884992	max_pooling3d_4[0][0]
conv3d_10 (Conv3D)	(None, 4, 4, 8, 256) 1769728	conv3d_9[0][0]
conv3d_transpose_1 (Conv3DTrans	(None, 8, 8, 16, 128 262272	conv3d_10[0]
conv3d_transpose_1 [0]		
concatenate_1 (Concatenate)	(None, 8, 8, 16, 256 0	conv3d_transpo
concatenate_1 [0][0]		conv3d_8[0][0]
conv3d_11 (Conv3D)	(None, 8, 8, 16, 128 884864	concatenate_1
conv3d_11 [0][0]		
conv3d_12 (Conv3D)	(None, 8, 8, 16, 128 442496	conv3d_11[0]
conv3d_12 [0]		
conv3d_transpose_2 (Conv3DTrans	(None, 16, 16, 32, 6 65600	conv3d_12[0]
conv3d_transpose_2 [0]		
concatenate_2 (Concatenate)	(None, 16, 16, 32, 1 0	conv3d_transpo
concatenate_2 [0][0]		conv3d_6[0][0]
conv3d_13 (Conv3D)	(None, 16, 16, 32, 6 221248	concatenate_2
conv3d_13 [0][0]		
conv3d_14 (Conv3D)	(None, 16, 16, 32, 6 110656	conv3d_13[0]
conv3d_14 [0]		
conv3d_transpose_3 (Conv3DTrans	(None, 32, 32, 64, 3 16416	conv3d_14[0]
conv3d_transpose_3 [0]		
concatenate_3 (Concatenate)	(None, 32, 32, 64, 6 0	conv3d_transpo
concatenate_3 [0][0]		conv3d_4[0][0]

conv3d_15 (Conv3D) [0][0]	(None, 32, 32, 64, 3 55328	concatenate_3
conv3d_16 (Conv3D) [0]	(None, 32, 32, 64, 3 27680	conv3d_15[0]
conv3d_transpose_4 (Conv3DTrans [0]	(None, 64, 64, 128, 4112	conv3d_16[0]
concatenate_4 (Concatenate) se_4[0][0]	(None, 64, 64, 128, 0	conv3d_transpo conv3d_2[0][0]
conv3d_17 (Conv3D) [0][0]	(None, 64, 64, 128, 13840	concatenate_4
conv3d_18 (Conv3D) [0]	(None, 64, 64, 128, 6928	conv3d_17[0]
conv3d_19 (Conv3D) [0]	(None, 64, 64, 128, 34	conv3d_18[0]
=====		
=====		
Total params: 5,644,930		
Trainable params: 5,644,930		
Non-trainable params: 0		

```
In [39]: # history = model.fit(x=X_train, y=y_train, batch_size=2, epochs=50, verbose=1)
# history = model.fit_generator(noaug_gen, steps_per_epoch=3, epochs=20, verbose=1)
checkpoint = ModelCheckpoint(filepath='weights3.h5')

history = model.fit_generator(train_gen, steps_per_epoch=30, epochs=20, verbose=1)
```

Epoch 1/20

2/30 [=>.....] - ETA: 1:16 - loss: 0.0189 - weighted_categorical_crossentropy: 0.0189

C:\Users\Gerald\Anaconda3\envs\tf3\lib\site-packages\scipy\ndimage\interpolation.py:600: UserWarning: From scipy 0.13.0, the output shape of zoom() is calculated with round() instead of int() - for these inputs the size of the returned array has changed.

"the returned array has changed.", UserWarning)

30/30 [=====] - 58s 2s/step - loss: 0.0211 - weighted_categorical_crossentropy: 0.0211

Epoch 2/20

30/30 [=====] - 54s 2s/step - loss: 0.0208 - weighted_categorical_crossentropy: 0.0208

Epoch 3/20

30/30 [=====] - 55s 2s/step - loss: 0.0190 - weighted_categorical_crossentropy: 0.0190

Epoch 4/20

30/30 [=====] - 56s 2s/step - loss: 0.0129 - weighted_categorical_crossentropy: 0.0129

Epoch 5/20

30/30 [=====] - 57s 2s/step - loss: 0.0209 - weighted_categorical_crossentropy: 0.0209

Epoch 6/20

30/30 [=====] - 56s 2s/step - loss: 0.0132 - weighted_categorical_crossentropy: 0.0132

Epoch 7/20

30/30 [=====] - 55s 2s/step - loss: 0.0075 - weighted_categorical_crossentropy: 0.0075

Epoch 8/20

30/30 [=====] - 55s 2s/step - loss: 0.0056 - weighted_categorical_crossentropy: 0.0056

Epoch 9/20

30/30 [=====] - 56s 2s/step - loss: 0.0047 - weighted_categorical_crossentropy: 0.0047

Epoch 10/20

30/30 [=====] - 55s 2s/step - loss: 0.0082 - weighted_categorical_crossentropy: 0.0082

Epoch 11/20

30/30 [=====] - 55s 2s/step - loss: 0.0055 - weighted_categorical_crossentropy: 0.0055

Epoch 12/20

30/30 [=====] - 55s 2s/step - loss: 0.0061 - weighted_categorical_crossentropy: 0.0061

Epoch 13/20

30/30 [=====] - 55s 2s/step - loss: 0.0042 - weighted_categorical_crossentropy: 0.0042


Epoch 14/20

30/30 [=====] - 56s 2s/step - loss: 0.0032 - weighted_

```
categorical_crossentropy: 0.0032
Epoch 15/20
30/30 [=====] - 56s 2s/step - loss: 0.0032 - weighted_
categorical_crossentropy: 0.0032
Epoch 16/20
30/30 [=====] - 56s 2s/step - loss: 0.0031 - weighted_
categorical_crossentropy: 0.0031
Epoch 17/20
30/30 [=====] - 56s 2s/step - loss: 0.0027 - weighted_
categorical_crossentropy: 0.0027
Epoch 18/20
30/30 [=====] - 55s 2s/step - loss: 0.0038 - weighted_
categorical_crossentropy: 0.0038
Epoch 19/20
30/30 [=====] - 55s 2s/step - loss: 0.0029 - weighted_
categorical_crossentropy: 0.0029
Epoch 20/20
30/30 [=====] - 55s 2s/step - loss: 0.0029 - weighted_
categorical_crossentropy: 0.0029
```

```
In [46]: plt.plot(history.history['loss'])
plt.ylabel('Training loss')
```

```
Out[46]: <matplotlib.text.Text at 0x1e099cb87f0>
```



```

In [62]: # make predictions on non-augmented data
noaug_gen.reset_index()
y_predict = model.predict_generator(noaug_gen)

# plot predictions
noaug_gen.reset_index()
X, y_true = noaug_gen.__getitem__(0)
y_pred=np.expand_dims(y_predict[:, :, :, :1], axis=4)
# plot the predictions
idx = [18, 19, 20, 27, 28]
plot_masks(X[2][idx], (y_true[2][idx]>1).astype('float'), y_pred[2][idx])

# save class labels
y_argmax = np.argmax(y_predict, axis=4).astype('float')
y_argmax *= 255/y_argmax.max()
y_argmax = np.round(y_argmax).astype('uint8')

# prediction probability
y_predict = y_predict[:, :, :, :1]*255
y_predict = y_predict.astype('uint8')
print(y_predict.mean())

# save prediction volumes to disk
for (i, vol) in enumerate(vols):
    print(vol)
    outfile = os.path.join(vol_folder, vol) + '_prob.tif'
    imsave(outfile, y_predict[i])

    outfile = os.path.join(vol_folder, vol) + '_pred.tif'
    imsave(outfile, y_argmax[i])

```

```

13.136070251464844
GP1825
GP1826
GP1827

```